

## Database Programming with SQL

### 9-1: Using GROUP BY and HAVING Clauses

- **HAVING** - Used to specify which groups are to be displayed; restricts groups that do not meet group criteria
- **GROUP BY** - Divides the rows in a table into groups

1. In the SQL query shown below, which of the following is true about this query?

```
SELECT last_name, MAX(salary)
FROM employees
WHERE last_name LIKE 'K%'
GROUP BY manager_id, last_name
HAVING MAX(salary) > 16000
ORDER BY last_name DESC ;
```

**c. Only salaries greater than 16001 will be in the result set.**

2. Each of the following SQL queries has an error. Find the error and correct it. Use Oracle Application Express to verify that your corrections produce the desired results.

a. SELECT manager\_id  
FROM employees  
WHERE AVG(salary) < 16000  
GROUP BY manager\_id;  
**SELECT manager\_id**  
**FROM employees**  
**GROUP BY manager\_id**  
**HAVING AVG(salary) < 16000;**

b. SELECT cd\_number, COUNT(title)  
FROM d\_cds  
WHERE cd\_number < 93;  
**SELECT cd\_number, COUNT(title)**  
**FROM d\_cds**  
**WHERE cd\_number < 93**  
**GROUP BY cd\_number;**

c. SELECT ID, MAX(ID), artist AS Artist  
FROM d\_songs  
WHERE duration IN('3 min', '6 min', '10 min')  
HAVING ID < 50  
GROUP by ID;

**SELECT ID, MAX(ID), artist**  
**FROM d\_songs**  
**WHERE duration IN ('3 min', '6 min', '10 min')**  
**GROUP BY ID, artist**  
**HAVING ID < 50;**

d. SELECT loc\_type, rental\_fee AS Fee  
FROM d\_venues  
WHERE id <100  
GROUP BY "Fee"  
ORDER BY 2;

**SELECT loc\_type, rental\_fee AS Fee**  
**FROM d\_venues**  
**WHERE id < 100**  
**GROUP BY loc\_type, rental\_fee**  
**ORDER BY 2;**

3. Rewrite the following query to accomplish the same result:

SELECT DISTINCT MAX(song\_id)  
FROM d\_track\_listings  
WHERE track IN ( 1, 2, 3);

**SELECT MAX(song\_id)**  
**FROM d\_track\_listings**  
**WHERE track IN (1, 2, 3);**

4. Indicate True or False

  T   a. If you include a group function and any other individual columns in a SELECT clause,

then each individual column must also appear in the GROUP BY clause.

\_\_F\_\_ b. You can use a column alias in the GROUP BY clause.

\_\_F\_\_ c. The GROUP BY clause always includes a group function.

5. Write a query that will return both the maximum and minimum average salary grouped by department from the employees table.

```
SELECT department_id, MAX(AVG(salary)) AS Max_Avg_Salary, MIN(AVG(salary))  
AS Min_Avg_Salary
```

```
FROM employees
```

```
GROUP BY department_id;
```

6. Write a query that will return the average of the maximum salaries in each department for the employees table.

```
SELECT AVG(MAX(salary)) AS Avg_Max_Salary
```

```
FROM employees
```

```
GROUP BY department_id;
```

## 9-2: Using ROLLUP and CUBE Operations and GROUPING SETS

- **ROLLUP** - Used to create subtotals that roll up from the most detailed level to a grand total, following a grouping list specified in the clause
- **CUBE** - An extension to the GROUP BY clause like ROLLUP that produces cross-tabulation reports
- **GROUPING SETS** - Used to specify multiple groupings of data

1. Within the Employees table, each manager\_id is the manager of one or more employees who each have a job\_id and earn a salary. For each manager, what is the total salary earned by all of the employees within each job\_id? Write a query to display the Manager\_id, job\_id, and total salary. Include in the result the subtotal salary for each manager and a grand total of all salaries.

```
SELECT manager_id, job_id, SUM(salary) AS total_salary
```

```
FROM employees
```

**GROUP BY ROLLUP(manager\_id, job\_id);**

2. Amend the previous query to also include a subtotal salary for each job\_id regardless of the manager\_id.

**SELECT job\_id, manager\_id, SUM(salary) AS total\_salary**

**FROM employees**

**GROUP BY ROLLUP(job\_id, manager\_id);**

3. Using GROUPING SETS, write a query to show the following groupings:

- department\_id, manager\_id, job\_id
- manager\_id, job\_id
- department\_id, manager\_id

**SELECT department\_id, manager\_id, job\_id, SUM(salary) AS total\_salary**

**FROM employees**

**GROUP BY GROUPING SETS (**

**(department\_id, manager\_id, job\_id),**

**(manager\_id, job\_id),**

**(department\_id, manager\_id)**

**);**

### 9-3: Set Operators

- **UNION** - operator that returns all rows from both tables and eliminates duplicates
- **NULL Columns** - columns that were made up to match queries in another table that are not in both tables
- **UNION ALL** - operator that returns all rows from both tables, including duplicates
- **SET Operators** - used to combine results into one single result from multiple SELECT statements
- **MINUS** - operator that returns rows that are unique to each table
- **INTERSECT** - operator that returns rows common to both tables

1. Name the different Set operators?

**UNION, UNION ALL, INTERSECT, MINUS (or EXCEPT)**

2. Write one query to return the employee\_id, job\_id, hire\_date, and department\_id of all employees and a second query listing employee\_id, job\_id, start\_date, and department\_id from the job\_history table and combine the results as one single output. Make sure you suppress duplicates in the output.

**SELECT employee\_id, job\_id, hire\_date AS date, department\_id**

**FROM employees**

**UNION**

**SELECT employee\_id, job\_id, start\_date AS date, department\_id**

**FROM job\_history;**

3. Amend the previous statement to not suppress duplicates and examine the output. How many extra rows did you get returned and which were they? Sort the output by employee\_id to make it easier to spot.

**SELECT employee\_id, job\_id, hire\_date AS date, department\_id**

**FROM employees**

**UNION ALL**

**SELECT employee\_id, job\_id, start\_date AS date, department\_id**

**FROM job\_history**

**ORDER BY employee\_id;**

4. List all employees who have not changed jobs even once. (Such employees are not found in the job\_history table)

**SELECT employee\_id**

**FROM employees**

**MINUS**

**SELECT employee\_id**

**FROM job\_history;**

5. List the employees that HAVE changed their jobs at least once.

**SELECT employee\_id**

**FROM employees**

**INTERSECT**

**SELECT employee\_id**

**FROM job\_history;**

6. Using the UNION operator, write a query that displays the employee\_id, job\_id, and salary of ALL present and past employees. If a salary is not found, then just display a 0 (zero) in its place.

**SELECT employee\_id, job\_id, COALESCE(salary, 0) AS salary**

**FROM employees**

**UNION**

**SELECT employee\_id, job\_id, COALESCE(salary, 0) AS salary**

**FROM job\_history;**